

Design and Implementation of a Workflow-Enabled Enterprise Document Management System with Legislative Compliance

Juraj Suchan

Department of Applied Informatics

Faculty of Informatics

Pan-European University

Field of Study: Informatics

Bratislava, Slovakia

Email: xsuchan@paneurouni.com

Index Terms—document management system, workflow engine, microservices, records management, legislative compliance

Abstract—This paper presents the design and implementation of an enterprise document and records management system with integrated workflow support. The proposed solution addresses the need for structured management of registry records, including classification, retention policies, disposal procedures, and auditability, while maintaining extensibility and modern architectural principles. The system is based on a microservice-oriented architecture implemented in .NET, utilizing a relational database for structured metadata storage and distributed components for messaging and caching.

A dedicated workflow engine enables the definition and execution of configurable process models, including task, decision, and scripted nodes, allowing dynamic business logic execution and versioned workflow definitions. The data model reflects hierarchical classification structures and regulatory requirements related to records lifecycle management. The implementation emphasizes modularity, traceability, and compliance with records management principles.

The resulting system demonstrates that enterprise-grade document and records management functionality can be combined with flexible workflow execution within a unified architecture, while preserving scalability and maintainability.

I. INTRODUCTION

Enterprise document and records management systems (DMS/RMS) play a critical role in ensuring structured storage, traceability, and lifecycle control of organizational documents. Modern organizations are required not only to manage digital documents efficiently, but also to comply with regulatory requirements concerning retention periods, classification schemes, disposal procedures, and auditability.

Traditional document management solutions often focus primarily on storage and retrieval, while advanced business process handling is delegated to separate workflow systems. This separation introduces integration complexity, data redundancy, and inconsistent lifecycle control. There is therefore a need for a unified approach that combines records management principles with flexible workflow execution in a single coherent architecture.

This paper presents the design and implementation of a prototype enterprise document and records management system with integrated workflow support. The system is built using modern architectural principles and emphasizes modularity, scalability, and compliance-oriented data modeling. The core contribution of this work lies in the integration of a configurable workflow engine with a structured records model reflecting hierarchical classification, retention rules, and audit requirements.

The objective of this paper is to describe the architectural design, data model, workflow engine structure, and compliance mechanisms of the proposed solution, and to demonstrate how these components form a consistent and extensible system suitable for enterprise environments.

II. BACKGROUND AND RELATED WORK

Document and records management systems have been extensively studied in the context of enterprise information systems. Standards such as MoReq and ISO 15489 define principles for records lifecycle management, including classification, retention policies, disposal procedures, and auditability. These frameworks emphasize the importance of structured metadata, hierarchical classification schemes, and controlled records disposition processes.

Modern enterprise software increasingly adopts service-oriented or microservice-based architectures to improve scalability and modularity. In parallel, workflow management systems have evolved to support dynamic process modeling, configurable execution logic, and distributed task orchestration. However, many existing solutions treat document management and workflow orchestration as separate subsystems, connected through integration layers.

This separation can lead to inconsistencies in state management, duplication of lifecycle information, and increased architectural complexity. Several research works propose tighter integration between content management and process execution, yet practical implementations often prioritize either document storage or workflow flexibility.

The solution presented in this paper builds upon established records management principles while integrating a configurable workflow engine directly into the core system architecture. The objective is not to redefine records management standards, but to demonstrate how compliance-oriented data modeling and workflow execution can coexist within a unified enterprise system.

III. SYSTEM ARCHITECTURE

The proposed system follows a modular service-oriented architecture designed to ensure scalability, extensibility, and separation of concerns. The backend is implemented using the .NET 10 platform and is structured into multiple logical components responsible for records management, workflow execution, authentication, and supporting infrastructure services.

Structured metadata related to records, classification hierarchies, workflow definitions, and execution states are stored in a relational database (Microsoft SQL Server). This choice enables strong consistency guarantees and transactional integrity for compliance-sensitive operations.

Binary document content is stored separately from structured metadata. Uploaded files are persisted in a document-oriented storage system (MongoDB), which is optimized for handling large binary objects and scalable file storage. This architectural separation between metadata and file content improves flexibility, supports horizontal scaling of storage components, and reduces performance bottlenecks in relational persistence layers.

Inter-service communication and asynchronous processing are supported by a message broker (RabbitMQ), which facilitates decoupled execution of background tasks and workflow-related events. A distributed caching layer (Redis) is employed to improve performance for frequently accessed metadata and configuration data.

Authentication and authorization are handled via an external identity provider (Keycloak), enabling standards-based authentication protocols and role-based access control. The system exposes its functionality through a REST-based API layer, allowing integration with external systems and frontend clients.

The user interface is implemented as a single-page application using React and TypeScript. A component-based UI framework is employed to ensure consistency and maintainability. This separation between backend services and frontend client supports independent evolution of presentation and business logic layers.

For deployment, the system is containerized and supports Docker-based orchestration. The architecture is designed to be environment-agnostic, enabling both local development deployments and cloud-ready configurations. This infrastructure approach aligns with modern enterprise deployment practices and ensures portability across different hosting environments.

IV. WORKFLOW ENGINE DESIGN

The workflow engine represents a central architectural component of the proposed system. Its purpose is to integrate

process orchestration directly into the records management platform, thereby eliminating the need for external workflow subsystems and reducing integration complexity. The engine is designed to support configurability, version control, persistent execution state, and dynamic evaluation of process logic.

Unlike traditional hard-coded business logic implementations, the workflow layer allows process models to be defined, modified, and versioned independently of application redeployment. This design enables organizations to adapt their document-related processes (e.g., approvals, registrations, reviews, or disposal procedures) without modifying the core application codebase.

A. Workflow Definition Model

Workflow processes are modeled as directed graphs consisting of nodes (steps) and transitions. The definition layer includes explicit support for versioning: each workflow definition is stored together with a snapshot of its structure, and newer versions maintain a reference to the previous version. This approach allows controlled evolution of workflow logic while preserving historical definitions for traceability and reproducibility.

The engine supports multiple node types, including start and end nodes, human task nodes, decision nodes, and script nodes. Transitions define the control flow between nodes and may include conditions or metadata used by the execution layer.

B. Execution Model

Workflow execution is represented by workflow instances that reference a specific version of a workflow definition. Each instance maintains persistent execution state, including the current node, completed steps, and contextual data required to evaluate transitions and decision logic. Persisting execution state enables fault tolerance and supports auditing of workflow progress over time.

Asynchronous execution and background processing are supported through message-based task dispatching, allowing workflow-related activities to be decoupled from synchronous API calls. This model improves scalability and helps ensure that long-running workflow steps do not block user interactions.

C. Script Nodes and Dynamic Logic

To enable flexible business logic without redeployment, the workflow engine provides script nodes. These nodes execute embedded scripts during runtime to evaluate conditions, compute values, or perform lightweight validation and routing logic. Script node outputs can influence transition selection and enable dynamic behavior in workflow graphs while keeping the core execution engine stable.

Overall, the workflow engine design balances configurability and maintainability by separating workflow structure (definitions) from runtime state (executions), while ensuring that versioned definitions remain available for auditing and consistent replay of historical workflow instances.

V. LEGISLATIVE COMPLIANCE MODEL

The records management component of the system is designed to reflect established principles of structured classification and lifecycle control. The data model incorporates hierarchical classification entities that represent subject-based groupings of records. These classification structures allow consistent assignment of documents to predefined categories and support standardized numbering schemes.

The core domain model includes entities representing evidence types, volumes, records, and case files. This structure enables separation between abstract record categories and their concrete yearly or contextual instances. Such modeling supports scalable growth of registry data while preserving logical grouping and traceability across different operational contexts.

Retention and disposal mechanisms are incorporated into the data model as lifecycle attributes associated with classification entities. While the current implementation provides structural support for defining retention periods and disposal types, the lifecycle automation layer is designed for further extensibility. This approach allows gradual refinement of compliance logic without altering the core data schema.

To ensure traceability and accountability, the system maintains audit-related metadata across key entities. Creation and modification timestamps, as well as user attribution fields, are persistently stored and integrated with the authentication subsystem. This design supports transparency and aligns with compliance-oriented requirements for records management systems.

By integrating hierarchical classification, lifecycle attributes, and auditability into the core domain model, the system establishes a structured foundation for regulatory compliance while remaining adaptable to evolving organizational requirements.

VI. DISCUSSION

The presented system demonstrates that integrating records management and workflow execution within a unified architecture can significantly reduce architectural fragmentation typically observed in enterprise environments. By avoiding separation between document storage and process orchestration layers, the system maintains consistent state management and minimizes synchronization complexity.

The service-oriented backend architecture enables horizontal scalability and modular evolution of individual components. The use of asynchronous messaging for workflow-related events decouples execution logic from synchronous API interactions, improving responsiveness and resilience under load. The separation between workflow definitions and runtime instances further enhances maintainability, as structural changes in process models do not affect already running workflow executions.

From a performance perspective, the introduction of a distributed caching layer reduces database load for frequently accessed configuration and metadata entities. The reliance on a relational database for compliance-critical operations ensures

transactional integrity and predictable behavior, which is particularly important in records lifecycle management scenarios.

Despite these advantages, certain limitations remain. While the system provides structural support for retention policies and disposal types, automated lifecycle enforcement mechanisms require further refinement and real-world validation. Additionally, large-scale deployment scenarios would benefit from extended performance benchmarking and distributed execution analysis.

Nevertheless, the architectural design establishes a solid technical foundation for enterprise-grade document and records management systems, combining configurability, traceability, and scalability within a cohesive framework.

VII. CONCLUSION

This paper presented the design and implementation of a prototype enterprise document and records management system with integrated workflow support. The proposed solution combines structured classification modeling, compliance-oriented data design, and a configurable workflow engine within a modular service-based architecture.

By separating workflow definitions from execution state, incorporating version control mechanisms, and integrating auditability across core entities, the system provides a flexible yet controlled environment suitable for enterprise document lifecycle management. The architecture leverages modern backend technologies and container-based deployment principles to ensure portability and scalability.

The results demonstrate that records management and workflow orchestration can be effectively unified without compromising architectural clarity or maintainability. Future work will focus on enhancing lifecycle automation mechanisms, extending compliance validation features, and evaluating performance characteristics under large-scale enterprise workloads.